

# Input-Output Organization: Programmed I/O, Interrupts, and Direct Memory Access

Based on Carl Hamacher, Computer Organization (6th Edition)

August 28, 2025

# Outline

1 Input-Output Organization

2 Program-Controlled I/O

3 Interrupts

# Introduction to I/O Organization

- The I/O subsystem connects the CPU and memory with the external environment, enabling interaction with devices like keyboards, monitors, disks, and networks.
- The main challenges addressed by I/O organization:
  - Speed mismatch between fast CPU and slow devices.
  - Different data formats between CPU (binary words) and devices (characters, signals).
  - Managing multiple devices simultaneously.
- Efficient I/O organization is critical for overall system performance.

# Types of I/O Devices

- Input Devices: keyboard, mouse, scanner – provide data to CPU.
- Output Devices: monitor, printer, speaker – receive processed results from CPU.
- Storage Devices: HDD, SSD, USB drives – allow long-term data storage.
- Communication Devices: Network Interface Cards (NICs), modems – connect system to networks.
- Each type of device requires specialized handling due to speed and data format differences.

# I/O Transfer Techniques

- Program-controlled I/O: CPU executes explicit instructions to monitor devices and transfer data.
- Interrupt-driven I/O: Device notifies CPU when ready, reducing CPU idle time.
- Direct Memory Access (DMA): Transfers large blocks directly between device and memory with minimal CPU intervention.

# Synchronous vs Asynchronous I/O

- **Synchronous I/O:**

- CPU waits for device to complete operation.
- Simple but inefficient if device is slow.

- **Asynchronous I/O:**

- CPU and device work independently.
  - Device signals CPU (via interrupt) when finished.
  - Increases CPU utilization.
- Handshaking protocols ensure correct timing between CPU and I/O device.

# I/O Interface Components

- **Data Register:** Holds data being transferred.
- **Status Register:** Indicates device state (ready, busy, error).
- **Control Register:** Used by CPU to issue commands to the device.
- Control logic translates CPU commands into signals devices can understand.
- Together, these registers form the I/O interface.

# Memory-Mapped vs I/O-Mapped I/O

- **I/O Mapped I/O:**

- Separate instructions for I/O (IN, OUT).
- Limited address space for devices.

- **Memory-Mapped I/O:**

- Device registers treated as memory addresses.
- Uniform instruction set for memory and I/O.
- Simplifies CPU design but reduces memory address space.

# Handshaking in I/O

- Handshaking ensures synchronization between CPU and slower devices.
- Example signals:
  - READY/BUSY – device signals readiness.
  - ACK/NACK – CPU acknowledges data.
- Prevents data loss and ensures proper timing.
- Example: A printer signals READY when buffer is available for next character.

# Typical I/O Transaction Cycle

- ① CPU writes command into device control register.
- ② Device executes operation and updates status register.
- ③ CPU polls status register or waits for interrupt.
- ④ Data transfer occurs via data register.

This sequence applies to both input and output devices.

# Problems with CPU-Controlled I/O

- CPU must frequently check device status.
- Leads to busy waiting and poor utilization.
- Cannot keep up with high-speed I/O (like disks or networks).
- Motivates use of interrupts and DMA.

## Section 8.1 Summary

- I/O subsystem bridges CPU and peripherals with major speed and format differences.
- Interfaces provide data, control, and status registers.
- Program-controlled I/O is inefficient, leading to advanced methods like interrupts and DMA.

# Concept of Program-Controlled I/O

- CPU takes full responsibility for monitoring device readiness.
- Data transfers are initiated only when device indicates ready in its status register.
- Example: Keyboard input – CPU checks the READY bit before reading data.

# Working of Programmed I/O

- The CPU executes instructions to manage I/O transfers step by step.
- Each transfer involves:
  - ① An **input instruction** to move data from device → CPU register.
  - ② A **store instruction** to move data from CPU register → memory.
- No direct path exists between the I/O device and main memory.

# CPU Involvement in Programmed I/O

- The CPU stays in a **program loop** (busy waiting).
- It repeatedly checks the device status register until the I/O device signals ready.
- This wastes valuable CPU cycles since no useful computation is done in the meantime.
- Example: Printer taking milliseconds while CPU operates in nanoseconds.

# Instruction Sequence Example

Example sequence for reading input:

- ① Read status register.
- ② If READY bit = 0, repeat.
- ③ If READY bit = 1, read data register.
- ④ Transfer data to memory or CPU register.

This cycle repeats for each word, making it inefficient for bulk transfers.

# Register Transfer Notation

- Example program fragment:
  - MOV R1, STATUS (load device status)
  - TEST R1, READY (check ready bit)
  - JZ WAIT (loop if not ready)
  - IN R2, DATA (read data)
  - STORE R2, MEM[addr] (write to memory)
- Every word transfer requires CPU intervention.

# Limitations of Program-Controlled I/O

- CPU utilization is extremely low.
- Only practical for very slow devices or occasional transfers.
- Cannot handle simultaneous I/O from multiple devices.
- Example: Network card generating thousands of packets per second.

# Flowchart Representation

Flow of program-controlled I/O:

- ① Poll device status.
- ② If not ready, repeat loop.
- ③ If ready, perform transfer.
- ④ Return to step 1 for next word.

Very inefficient but simple to implement.

# Performance Issues

- Modern CPUs operate at GHz, I/O devices at kHz–MHz.
- Polling wastes millions of cycles.
- Performance gap widens with faster CPUs.
- Necessitates more advanced approaches.

## Section 8.2 Summary

Program-controlled I/O is simple but leads to poor CPU efficiency. Better methods are interrupts and DMA.

# Introduction to Interrupts

- Interrupts allow devices to signal CPU when they are ready.
- CPU can execute other tasks instead of waiting.
- Provides asynchronous handling of I/O events.

# Interrupt-Initiated I/O: Concept

- Unlike programmed I/O, the CPU does **not** constantly poll for device readiness.
- The I/O interface itself monitors peripherals and sends an **interrupt request** when data is ready.
- The CPU continues executing other tasks until interrupted.

# Sequence of Interrupt-I/O Operation

- ① CPU issues a command to prepare for I/O and then proceeds with other tasks.
- ② The I/O interface continuously monitors device status.
- ③ When the device is ready, the interface sends an **interrupt request (IRQ)**.
- ④ CPU suspends current execution, jumps to an Interrupt Service Routine (ISR).
- ⑤ ISR performs the I/O transfer and then control returns to the original program.

# Advantages of Interrupt-I/O

- CPU doesn't waste cycles polling devices.
- Improves system efficiency and responsiveness.
- Enables better overlap of computation and I/O.

# Limitations of Interrupt-I/O

- Transfer rate limited by CPU's ability to service interrupts.
- Each I/O operation still requires CPU involvement through ISR.
- Not as efficient as DMA for large block transfers.

# Types of Interrupts (Part 1)

- Hardware Interrupts: Generated by I/O devices through hardware lines.
- Software Interrupts: Triggered by instructions within the program.
- Vectored Interrupts: Predefined address of ISR.
- Non-vectored Interrupts: CPU must determine ISR dynamically.

## Types of Interrupts (Part 2)

- Maskable Interrupts: Can be enabled or disabled by the CPU.
- Non-maskable Interrupts (NMI): Cannot be disabled; highest priority.
- External Interrupts: Generated by hardware devices.
- Internal Interrupts: Generated internally by CPU (faults, exceptions).
- Synchronous vs Asynchronous Interrupts: Predictable (timer) vs unpredictable (I/O).

## Section 8.3 Summary

- Interrupts eliminate busy waiting and improve CPU utilization.
- Interrupt types allow flexible handling of I/O and system events.
- For high-speed bulk data, DMA is more efficient than interrupt-driven I/O.